

AI-102 Practice Assessment 1st Attempt

Question 1 of 50

You have a website that allows users to upload images.

You need to ensure that the uploaded images do not contain adult content. The solution must minimize development effort.

Which service should you use?

Your Answer

☒ Azure AI Vision Image Analysis

✓ This answer is correct.

Correct Answer

☒ Azure AI Vision Image Analysis

✓ This answer is correct.

The Azure AI Vision Image Analysis service can extract a wide variety of visual features from an image. One of them is to detect adult content.

The Azure AI Face service provides AI algorithms that detect, recognize, and analyze human faces in images. Azure AI Custom Vision is an image recognition service that lets you build, deploy, and improve own image identifier models. So, while it is possible, it is not the solution with the lowest development effort. Azure AI Vision Spatial Analysis is used to ingest streaming video from cameras, extract insights, and generate events to be used by other systems.

What are Azure Cognitive Services? - Azure Cognitive Services | Microsoft Learn

What is Computer Vision? - Azure Cognitive Services | Microsoft Learn

What is Image Analysis? - Azure Cognitive Services | Microsoft Learn

Plan and prepare to develop AI solutions on Azure

Question 2 of 50

You plan to implement a solution to extract information from documents by using Azure AI Document Intelligence and use prebuilt models where possible.

You need to identify the prebuilt models available in Azure AI Document Intelligence.

Which three models are available? Each correct answer presents a complete solution.

Your Answer

☒ invoice

✓ This answer is correct.

☒ receipt

✓ This answer is correct.

☒ W-2

✓ This answer is correct.

Question 3 of 50

You are building a knowledge mining solution that uses Azure AI Search.

You need to extract content from a file within the enrichment pipeline by using AI enrichment.

Which built-in skill should you use?

- ☐ `Microsoft.Skills.Text.KeyPhraseExtractionSkill`
- ☐ `Microsoft.Skills.Text.SplitSkill`
- ☐ `Microsoft.Skills.Text.V3.EntityRecognitionSkill`
- ☒ `Microsoft.Skills.Util.DocumentExtractionSkill`

✓ This answer is correct.

`Microsoft.Skills.Util.DocumentExtractionSkill` is the built-in skill used to extract content from a file within the enrichment pipeline.

[Built-in skills – Azure Cognitive Search | Microsoft Learn](#)

[Create an Azure Cognitive Search solution – Training | Microsoft Learn](#)

Your organization processes a large volume of scanned vendor documents, some of which are poorly scanned or contain creases.

You need to extract structured data such as vendor details, totals, and line items from these invoices.

What should you use to achieve this goal?

- ☐ Train a custom model in Azure AI Document Intelligence Studio.
- ☐ Use Azure AI Vision OCR.
- ☒ Use the Invoice model in Azure AI Document Intelligence.
- ☐ Use the Business Card model in Azure AI Document Intelligence.

✓ This answer is correct.

The Invoice model in Azure AI Document Intelligence is the optimal choice for processing invoices due to its specialized capabilities for extracting structured data, even from poorly scanned documents. Training a custom model is unnecessary because the Invoice model already meets the requirements. Azure AI Vision OCR lacks the ability to extract structured data, making it unsuitable for this task. The Business Card model is designed for business cards and does not provide the necessary functionality for processing invoices.

Use financial, ID, and tax models - [Training](#) | Microsoft Learn

Your company receives incoming documents that need to be classified into predefined categories before extracting specific data fields.

You need to implement a solution that accurately classifies documents and supports subsequent data extraction.

Which solution should you recommend?

- ☐ Azure AI Vision OCR (Optical Character Recognition)
- ☒ Custom classification model in Azure AI Document Intelligence
 - ✓ This answer is correct.
- ☐ Custom template model in Azure AI Document Intelligence
- ☐ Invoice model in Azure AI Document Intelligence

The custom classification model in Azure AI Document Intelligence is the most suitable solution because it is specifically designed to categorize documents into predefined categories, enabling accurate classification and supporting subsequent data extraction. Custom template models are limited to data extraction from structured documents and do not address classification needs. Azure AI Vision OCR focuses solely on text recognition and lacks the functionality required for document categorization. The Invoice model is tailored for invoice processing and does not meet the broader requirements for document classification described in the scenario.

[What is Azure Document Intelligence? - Training | Microsoft Learn](#)

[Use financial, ID, and tax models - Training | Microsoft Learn](#)

You have an app named App1 that extracts invoice data from PDF files by using an S0 instance of Azure AI Document Intelligence. The PDF files are up to 2 MB each and contain up to 10 pages.

Users report that App1 is unable to process some invoices.

You need to troubleshoot the issue.

What is a possible cause of the issue?

- ☒ Some of the files are password protected.
✓ This answer is correct.
- ☐ Some of the files are too large.
- ☐ Some of the files have too many pages.
- ☐ The tier of the Azure AI Document Intelligence instance is insufficient.

This answer is incorrect.

The service cannot process password-protected files, and this can cause the service a processing failure for some of the files. Although file size and number of pages can cause failures, the limit for the S0 tier is 500 MB and 2,000 pages.

The S0 tier is sufficient for the file characteristics mentioned.

[Invoice data extraction – Form Recognizer - Azure Applied AI Services | Microsoft Learn](#)

[Extract data from forms with Form Recognizer - Training | Microsoft Learn](#)

You are building an app that uses Azure AI Services Document Translation.

You need to improve the quality of the translation for user-uploaded documents.

What should you ask the users to include when they upload a document?

- ☐ a summary
- ☐ the file format
- ☒ the source language
- ☐ the writing style

✓ This answer is correct.

If the language of the content in the source document is known, it is recommended to specify the source language in the request to get a better translation.

[Frequently asked questions - Document Translation - Azure Cognitive Services | Microsoft Learn](#)

[Translate text with the Translator service - Training | Microsoft Learn](#)

Your organization processes invoices from vendors in various formats, including scanned PDFs and digital images.

You need to implement Azure AI Document Intelligence to extract key fields such as invoice number, due date, and total amount.

Each correct answer presents part of the solution. Which two actions should you perform?

- ☐ Create a custom model for invoice processing.
- ☐ Install the solution on a local server.
- ☐ Label all invoice documents manually.
- ☒ Provision an Azure AI Document Intelligence resource.
✓ This answer is correct.
- ☐ Edit the invoices to ensure consistency.

- ☒ Use the prebuilt invoice model to extract key fields.
✓ This answer is correct.

To implement the solution, provisioning the Azure AI Document Intelligence resource is a necessary first step to enable its capabilities. Using the prebuilt invoice model is the most efficient approach for extracting key fields such as invoice number, due date, and total amount, as it is specifically designed for this purpose. Installing the solution on a local server is irrelevant because the service is cloud-based and does not require local installation. Creating a custom model is unnecessary because the prebuilt invoice model already meets the requirements, saving time and effort. Manually labeling invoice documents is also not needed, as the prebuilt model leverages pre-trained capabilities to eliminate this step. Editing the invoices to ensure consistency won't help when processing incoming invoices and doesn't provide the model with varied training data

[Document Intelligence invoice model - Training | Microsoft Learn](#)

An organization manually processes a high volume of invoices daily, leading to errors in data extraction.

You need to use Microsoft Azure AI Document Intelligence to automate invoice data extraction and integration.

Each correct answer presents part of the solution. Which two actions should you take?

☒ Integrate the extracted data with an organization's database.

✓ This answer is correct.

☐ Train a custom model for document processing.

☐ Use a general-purpose model for text extraction.

☒ Use the prebuilt Invoice model for data extraction.

✓ This answer is correct.

Integrating the extracted data with an organization's database ensures the information is stored and used effectively within existing systems, supporting downstream processes and workflows. Using the prebuilt Invoice model for data extraction is ideal because it is specifically designed to handle invoice documents, extracting structured data such as key-value pairs and tables efficiently. Training a custom model for document processing is unnecessary in this scenario since the prebuilt Invoice model meets the requirements for standard invoices. Using a general-purpose model for text extraction does not align with the goal of extracting structured data from invoices, as it lacks the specificity required for this task.

[What is Azure Document Intelligence? - Training | Microsoft Learn](#)

[Understand AI Document Intelligence - Training | Microsoft Learn](#)

You are provisioning an Azure OpenAI in Foundry Models service resource.

You need to ensure that the resource is only available to applications that are hosted in your Azure subscription.

Which network security setting should you configure?

- ☐ All networks
- ☐ All networks, and a network security group to control traffic
- ☒ Disabled, and allow a private endpoint connection to establish access
- ☐ Selected networks

✓ This answer is correct.

Because the requirements state that access to the resource should only be for applications hosted in the Azure subscription, setting the network option to Disabled and configuring a private endpoint meets this requirement. All networks permit access to any network, including the internet. Selected networks access to networks outside of Azure, if configured that way.

[How-to: Create and deploy an Azure OpenAI Service resource - Azure OpenAI | Microsoft Learn](#)

You are building a GPT-based chat application that will answer questions about your company.

You plan to use the Using your data feature in Azure OpenAI to ground the model with your company data.

While testing, you discover that some responses are not accurate enough.

You need to configure the Azure OpenAI resource to filter out less-relevant documents for responses.

Which parameter should you configure?

☐ Content data

☐ File name

☒ Retrieved documents

This answer is incorrect.

☐ Strictness

✓ This answer is correct.

The Strictness parameter sets the threshold to categorize documents as relevant to your queries. Raising the Strictness parameter value means a higher threshold for relevance and filters out more less-relevant documents for responses. Retrieved documents specifies the number of top-scoring documents from your data index used to generate responses. Content data specifies the fields in your index that contain the main text content of each document. File name specifies the field in your index that contains the original file name of each document.

[Using your data with Azure OpenAI Service - Azure OpenAI | Microsoft Learn](#)

You are building a GPT-based chat application that will answer questions about your company.

You plan to test the application by using strategies defined by Microsoft best practices.

Which three prompt engineering strategies should you consider while testing the application? Each correct answer presents a complete solution.

☒ Be Descriptive

✓This answer is correct.

☐ Be minimalistic

☐ Be simple

☒ Be Specific

✓This answer is correct.

☒ Order Matters

✓This answer is correct.

Be Specific means to leave as little to interpretation as possible. Be Descriptive means to use analogies. Order Matters means that the order in which you present information to the model can affect the output. Therefore, those three are valid best practices. Be simple and be minimalistic do not produce the best results and are, therefore, not best practices

[Azure OpenAI Service - Azure OpenAI | Microsoft Learn](#)

You are creating an application that will use Azure OpenAI REST API services. The application uses a REST call to a DALL-E model to generate images. The three parameters in the REST call are `prompt`, `n`, and `size`.

What does the `size` parameter indicate?

☐ the number of responses that you want returned

☒ the size of the images in bytes

This answer is incorrect.

☐ the size of the images in kilobytes

☐ the size of the images in pixel resolution

✓ This answer is correct.

To make a REST call to the services, you need the endpoint and authorization key for the Azure OpenAI service resource you provisioned in Azure. You initiate the image generation process by submitting a `POST` request to the service endpoint that has the authorization key in the header. The request must contain the following parameters in a JSON body:

`prompt`: The description of the image to be generated

`n`: The number of images to be generated

`size`: The resolution of the image to be generated (256x256, 512x512, or 1024x1024)

Use the Azure OpenAI REST API to consume DALL-E models - Training | Microsoft Learn

Your organization is developing an AI-powered application using Azure OpenAI in Foundry Models.

You need to choose a model for text generation.

What should you recommend?

- ☐ Azure AI Vision
- ☐ Code-Davinci-003
- ☐ DALL-E

☒ GPT-4

✓ This answer is correct.

Azure AI Vision is designed for analyzing visual content and does not support text generation, making it unsuitable for the scenario. Code-Davinci-003 is optimized for code generation tasks and lacks the capabilities required for text generation, which disqualifies it as a viable option. DALL-E specializes in generating images from textual descriptions and does not meet the requirement for text generation. GPT-4, on the other hand, is specifically designed for generating detailed and contextually accurate text responses, making it the most appropriate choice for the given requirement.

[Azure AI services - Training | Microsoft Learn](#)

[Plan and prepare to develop AI solutions on Azure - Training | Microsoft Learn](#)

Your company is integrating Azure OpenAI in Foundry Models into its internal knowledge base system to provide intelligent query responses.

You need to ensure that the most relevant documents are prioritized during data retrieval.

What should you do?

- ☐ Enable hybrid search mode.
- ☐ Increase retrieved document count.
- ☐ Reduce chunk size.

☒ Set strictness to a higher value.

✓ This answer is correct.

Setting strictness to a higher value ensures that only the most relevant documents are retrieved, effectively prioritizing relevance during data retrieval. Enabling hybrid search mode combines different search methods but does not specifically focus on document relevance. Increasing the number of retrieved documents may include less relevant data, which contradicts the goal of prioritizing relevance. Reducing chunk size impacts data granularity but does not directly influence the prioritization of document relevance.

[Azure OpenAI On Your Data - Training | Microsoft Learn](#)

Your organization is developing a customer-facing application that uses Azure OpenAI in Foundry Models to generate personalized responses. The application is connected to a Microsoft Azure AI Search index.

You need to configure the application to retrieve relevant data from the search index.

What should you do?

☐ Deploy an additional Azure OpenAI model.

☒ Enable semantic search.

✓ This answer is correct.

☐ Increase chunk size for data ingestion.

☐ Use keyword search.

Enabling semantic search enhances the precision and relevance of search results by interpreting the meaning behind query terms, making it the most suitable solution for improving response accuracy in this scenario. Increasing chunk size impacts data processing but does not directly affect the accuracy of retrieved data. Deploying an additional Azure OpenAI model does not address the specific need for improving data retrieval accuracy from the search index. Using keyword search lacks semantic capabilities, which are essential for achieving the desired level of accuracy.

[Make your data searchable - Training | Microsoft Learn](#)

An organization plans to use Azure OpenAI in Foundry Models to process user input and generate content.

You need to create visual content from text descriptions.

What should you use?

☐ Azure AI Search

☒ DALL-E

✓ This answer is correct.

☐ GPT-4

☐ Text Embedding Models

DALL-E is the appropriate choice for creating visual content from text descriptions, as it is specifically designed for this purpose. Azure Cognitive Search focuses on information retrieval and does not support visual content creation. GPT-4 is optimized for text-based tasks and lacks the capability to generate visual content. Text Embedding Models are intended for embedding tasks and are unrelated to the goal of creating visual content from text descriptions.

[Plan and prepare to develop AI solutions on Azure - Training | Microsoft Learn](#)

[Developer tools and SDKs - Training | Microsoft Learn](#)

Your company uses Azure OpenAI in Foundry Models to analyze large datasets.

You need to configure the application to retrieve relevant data efficiently.

What action should you take?

☐ Deploy a GPT-4 model.

☒ Enable vector search.

✓ This answer is correct.

☐ Increase chunk size.

☒ Use keyword-based search.

This answer is incorrect.

Enabling vector search enhances the relevance and efficiency of data retrieval, meeting the stated goal. Deploying an additional model does not address the need for efficient data retrieval, making it unsuitable. Increasing chunk size affects data processing but does not directly improve retrieval efficiency. Using keyword-based search lacks semantic capabilities, which are critical for achieving the desired outcome.

[Make your data searchable - Training | Microsoft Learn](#)

[Azure OpenAI On Your Data - Training | Microsoft Learn](#)

You have a website that allows users to upload images.

You need to ensure that the uploaded images do not contain adult content. The solution must minimize development effort.

Which service should you use?

- ☐ Azure AI Face Service
- ☐ Azure AI Custom Vision
- ☒ Azure AI Vision Image Analysis
- ☐ Azure AI Vision Spatial Analysis

✓ This answer is correct.

The Azure AI Vision Image Analysis service can extract a wide variety of visual features from an image. One of them is to detect adult content.

The Azure AI Face service provides AI algorithms that detect, recognize, and analyze human faces in images. Azure AI Custom Vision is an image recognition service that lets you build, deploy, and improve own image identifier models. So, while it is possible, it is not the solution with the lowest development effort. Azure AI Vision Spatial Analysis is used to ingest streaming video from cameras, extract insights, and generate events to be used by other systems.

[What are Azure Cognitive Services? - Azure Cognitive Services | Microsoft Learn](#)

[What is Computer Vision? - Azure Cognitive Services | Microsoft Learn](#)

[What is Image Analysis? - Azure Cognitive Services | Microsoft Learn](#)

[Plan and prepare to develop AI solutions on Azure](#)

You plan to implement a solution to extract information from documents by using Azure AI Document Intelligence and use prebuilt models where possible.

You need to identify the prebuilt models available in Azure AI Document Intelligence.

Which three models are available? Each correct answer presents a complete solution.

☒ invoice

✓ This answer is correct.

☐ meeting minutes

☒ receipt

✓ This answer is correct.

☐ time sheets

☒ W-2

✓ This answer is correct.

Except for meeting minutes and time sheets, all the other models are prebuilt models in Azure AI Document Intelligence. Meeting minutes and time sheets must be added to Azure AI Document Intelligence as custom models.
[Use financial, ID, and tax models - Training | Microsoft Learn](#)

You are building a mobile app that will enable users to scan street signs and will read out the text on the sign.

You need to recommend a service to use. The solution must minimize development effort.

Which service should you recommend?

☐ Azure AI Vision

✓ This answer is correct.

☒ Azure AI Custom Vision

This answer is incorrect.

☐ Azure AI Face Service

☐ Azure AI Document Intelligence

Azure AI Vision is the only service which can achieve the desired result.

Azure AI Custom Vision and Azure AI Face do not offer OCR. Azure AI Document Intelligence is designed for documents, but not images.

[OCR - Optical Character Recognition - Azure Cognitive Services | Microsoft Learn](#)

[Read Text in Images and Documents with the Computer Vision Service - Training | Microsoft Learn](#)

You are building an Azure AI Language solution.

You need to deploy the solution to a location without internet connectivity.

What should you do?

☐ Deploy the solution to a Docker host container.

✓ This answer is correct.

☒ Download the model and deploy it to a virtual machine.

This answer is incorrect.

☐ Use an Azure AI Foundry Service standard instance.

You can use disconnected containers to host Language Understanding on-premises.

There is no virtual machine set up to run a Language Understanding instance and a standard instance still has Language Understanding running in the cloud.

[Use Docker containers in disconnected environments - Azure Cognitive Services | Microsoft Learn](#)

[Deploy cognitive services in containers - Training | Microsoft Learn](#)

You have an Azure AI Foundry Service resource.

You need to enable diagnostic logging.

What are two prerequisites for diagnostic logging? Each correct answer presents a complete solution.

☒ a Log Analytics workspace

✓This answer is correct.

☐ an Azure Cosmos DB for NoSQL account

☐ an Azure key vault

☐ an Azure SQL database

☒ an Azure Storage account

✓This answer is correct.

The prerequisites to enable diagnostic logging are to have an Azure Storage resource that retains diagnostic logs for policy audit, static analysis, or backup. A Log Analytics resource provides a flexible log search and analytics tool that allows for analysis of raw logs generated by an Azure resource.

[Diagnostic logging - Azure Cognitive Services | Microsoft Learn](#)

[Monitor Cognitive Services - Training | Microsoft Learn](#)

You are developing a containerized optical character recognition (OCR)-capable application by using Azure AI Services containers. While developing the solution, you retrieve a status message of "Mismatch", and the connection to the AI Services resource fails. You need to ensure that the solution can connect to the AI Services resource. What should you do?

- ☐ Confirm that the API key is for the correct region.
- ☒ Confirm that the API key is for the correct resource type.
- ☐ Confirm that the Azure AI Services resource is online.
- ☐ Upgrade the Azure AI Services resource to a higher tier.

✓ This answer is correct.

Mismatch means that the wrong API key has been used. If you have provided an API key or endpoint for a different kind of Azure AI Services resource, you find your API key and service region in the Azure portal, in the Keys and Endpoint section for your Azure AI Services resource.

If the API key is invalid, you must confirm that the API key is for the correct region. If the API key has exceeded the quota, then you can either upgrade your pricing tier or wait for an additional quota to become available. Find your tier in the Azure portal, in the Pricing Tier section of your Azure AI Service resource. The mismatch error would not be generated if the resource was offline.

[Cognitive Services containers FAQ - Azure Cognitive Services | Microsoft Learn](#)

[Deploy cognitive services in containers - Training | Microsoft Learn](#)

Question 25 of 50

You have a website that enables users to upload and share PNG image files.

You implement an Azure AI Foundry Content Safety solution for the website.

The users report that some images fail to display after they are uploaded.

You need to ensure that the solution assigns proper content flags to the uploaded images.

What should you do?

☒ Adjust the content categories to be used.

This answer is incorrect.

☐ Adjust the severity level.

✓ This answer is correct.

☐ Convert the images to a different file format.

☐ Remove the severity level.

Best practices for improving system performance:

Monitor the system's performance regularly to ensure that the tradeoff is appropriate for your use case.

Modify the severity levels for blocking based on user feedback and observed trends in content safety.

[Transparency Note and use cases for AI Content Safety - Azure AI services | Microsoft Learn](#)

You create an Azure AI Foundry Content Safety solution to monitor the text input of a chat forum.

You discover that the solution returns a severity level of 0 for various instances of harmful text.

Which type of accuracy error is returned by the solution?

☐ a false negative

✓ This answer is correct.

☒ a false positive

This answer is incorrect.

☐ a true negative

☐ a true positive

A false negative arises when the model fails to identify harmful content and returns a severity level of 0. There are four accuracy error types, and a false negative indicates incorrectly accepted content in the text.

[Transparency Note and use cases for AI Content Safety - Azure AI services | Microsoft Learn](#)

Question 27 of 50

You plan to build an AI solution that will detect harmful content in company documents.

Which Azure AI Foundry Service should you include in the solution?

☐ Azure AI Search

☐ Bot Service

☒ Content Safety

✓ This answer is correct.

☐ Document Intelligence

☐ Metrics Advisor

Content Safety is the only listed service that provides the ability to automatically flag content in text or images as harmful. Content Safety recognizes four distinct categories of objectionable content: Hate and fairness, sexual, violence, and self-harm.

[Harm categories in Azure AI Content Safety - Azure AI services | Microsoft Learn](#)

[What are Azure AI services? - Azure AI services | Microsoft Learn](#)

You are building a custom copilot based on Azure OpenAI.

You discover attempts to exploit the copilot by jailbreaking it through a User Prompt Injection Attack (UPIA).

You need to increase security to prevent the jailbreak attack.

Which Azure AI Foundry Service should you include in the solution?

☒ Content Safety

✓ This answer is correct.

☐ Face

☐ Language

☐ Video Indexer

☐ Vision

Content Safety jailbreak risk detection recognizes four different classes of jailbreak attacks: Attempt to change system rules, embedding a conversation mockup to confuse the model, role-play, and encoding attacks. These can be used in LLM-based applications to prevent jailbreak attacks. All the other services listed are designed for different purposes.

[Jailbreak risk detection in Azure AI Content Safety - Azure AI services | Microsoft Learn](#)

[What are Azure AI services? - Azure AI services | Microsoft Learn](#)

Your organization is implementing Azure AI Foundry Content Safety to moderate user-generated content on a social messaging platform.

You need to configure Foundry Content Safety to identify and restrict harmful text content, including profanities and hate speech, while allowing customization for specific patterns.

Each correct answer presents part of the solution. Which two actions should you take?

☒ Configure harmful patterns using the Custom Categories API.

✓ This answer is correct.

☒ Enable built-in blocklists in Content Safety Studio.

✓ This answer is correct.

☐ Enable encryption at rest with customer-managed keys.

☐ Use the Analyze Image API for text analysis.

Enabling built-in blocklists in Content Safety Studio is essential as it provides predefined terms to effectively flag harmful content such as profanities and hate speech. Configuring harmful patterns using the Custom Categories API allows for tailored moderation by defining specific patterns that align with organizational needs. Enabling encryption at rest with customer-managed keys enhances data security but does not contribute to the detection or restriction of harmful text content. Using the Analyze Image API focuses on image moderation and is not applicable to the task of moderating text content.

[Azure AI services - Training | Microsoft Learn](#)

[Use the Azure Document Intelligence Studio - Training | Microsoft Learn](#)

Your team is deploying an AI agent to perform actions based on user requests, such as scheduling meetings and sending notifications.

You need to integrate the AI agent with tools that enable it to perform these actions programmatically.

What should you do?

- ☒ Add custom functions.
✓ This answer is correct.
- ☐ Configure static templates.
- ☐ Enable default model capabilities.
- ☐ Use a pre-trained chatbot framework.

Adding custom functions enables the AI agent to dynamically execute specific actions such as scheduling meetings and sending notifications, making it the most suitable choice for the scenario. Configuring static templates does not provide the flexibility needed for dynamic action execution, limiting its effectiveness. Default model capabilities lack the ability to perform specific programmatic actions, making them insufficient for the task. Using a pre-trained chatbot framework without customization fails to address the specific requirements for executing actions programmatically, reducing its applicability in this context.

[Integrate custom tools into your agent - Training | Microsoft Learn](#)

Your company plans to deploy an AI agent to analyze customer feedback and generate insights from text.

You need to select the most suitable Azure AI Foundry Service for text analysis.

What should you recommend?

☒ Azure AI Language.

✓ This answer is correct.

☐ Azure AI Speech.

☐ Azure AI Translator.

☐ Azure AI Vision.

Azure AI Language is the most appropriate service for this scenario as it specializes in analyzing and deriving insights from natural language text, directly addressing the requirements for text analysis. Azure AI Speech is not suitable because it focuses on audio transformations such as speech-to-text and text-to-speech, which are unrelated to the task. Azure AI Translator is designed for translating text between languages and does not provide the necessary capabilities for analyzing and generating insights from text. Azure AI Vision is focused on image-related tasks and lacks the functionality required for processing natural language text.

[Plan and prepare to develop AI solutions on Azure - Training | Microsoft Learn](#)

[Azure AI services - Training | Microsoft Learn](#)

Your organization is developing an AI agent for customer support inquiries that must integrate with Azure resources.

You need to select a framework that supports AI orchestration and multi-agent workflows.

What should you recommend?

- ☐ Azure Bot Framework
- ☐ Azure Machine Learning Studio
- ☐ Cognitive Services API

☒ Semantic Kernel framework

✓ This answer is correct.

Azure Bot Framework is commonly used for building conversational bots but does not provide the orchestration features required for multi-agent workflows. Azure Machine Learning Studio is focused on developing machine learning models and does not address the need for orchestration in this scenario. Cognitive Services API offers pre-built AI functionalities but lacks the ability to orchestrate multi-agent workflows. Semantic Kernel framework is the most suitable choice as it supports generative AI orchestration and multi-agent workflows, aligning perfectly with the requirements of the scenario.

[Provision an Azure AI Vision resource - Training | Microsoft Learn](#)

[Developer tools and SDKs - Training | Microsoft Learn](#)

You are developing a custom AI agent using Azure AI Foundry Agent Service. The agent needs to interact with an Azure AI Foundry resource that has the GPT-35-Turbo model deployed.

You need to configure the agent to effectively use the GPT-35-Turbo model for generating responses.

What should you do?

- ☐ Deploy a GPT-4 model.
- ☒ Specify the GPT-35-Turbo deployment name.

✓ This answer is correct.

- ☐ Use the Completion endpoint.
- ☐ Use the Embeddings endpoint.

To configure the agent to effectively use the GPT-35-Turbo model, specifying the deployment name ensures the agent interacts with the correct model deployed in a Microsoft Azure OpenAI resource. Using the Completion endpoint is not suitable for this model, as the ChatCompletion endpoint is preferred, and these endpoints aren't applicable to agents. Deploying a GPT-4 model does not meet the requirement to use the GPT-35-Turbo model specified in the scenario. The Embeddings endpoint is not appropriate for generating responses, as it serves a different purpose, such as creating vector representations of text.

[Get started with AI agent development on Azure - Training | Microsoft Learn](#)

You are developing an AI agent using Azure AI Foundry Agent Service to assist with data retrieval and analysis tasks.

You need to configure the AI agent to interact effectively with external data sources.

What should you use?

☒ Configure Azure AI Search.

✓ This answer is correct.

☐ Deploy a custom model in Azure Machine Learning.

☐ Implement Azure Bot Framework.

☐ Use Azure AI Document Intelligence.

Configuring Azure AI Search is the most appropriate solution because it provides indexing and querying capabilities that enable effective interaction with external data sources. Deploying a custom model in Azure Machine Learning is not suitable as it focuses on model deployment rather than facilitating external data source interaction. Using Azure AI Document Intelligence is not appropriate because it is designed for extracting structured data from documents, which does not align with the broader goal of data retrieval and analysis. Implementing Azure Bot Framework is ineffective in this context as it is intended for building conversational interfaces and does not address the need for data retrieval and analysis.

[Developer tools and SDKs - Training | Microsoft Learn](#)

Question 35 of 50

You are building an app that will use Azure AI Vision to detect the presence of people in a video feed.

Which Azure AI Vision feature should you use?

☒ face detection

This answer is incorrect.

☐ Image Analysis

☐ optical character recognition (OCR)

☒ Spatial Analysis

✓ This answer is correct.

The only visual feature that provides this capability is Spatial Analysis, as OCR, Image Analysis, and face detection are not meant to analyze the presence of people in a video feed.

[What is Computer Vision? - Azure Cognitive Services | Microsoft Learn](#)

[Analyze video - Training | Microsoft Learn](#)

Question 36 of 50

You are building an app that will use Azure AI Vision to extract text from scanned images of handwritten text.

Which Azure AI Vision feature should you use?

☐ Azure AI Custom Vision

☐ handwriting analysis

☐ Image Analysis

☒ optical character recognition (OCR)

✓ This answer is correct.

☐ Spatial Analysis

OCR is the only visual feature that can extract text from images.

[What is Computer Vision? - Azure Cognitive Services | Microsoft Learn](#)

[Read Text in Images and Documents with the Computer Vision Service - Training | Microsoft Learn](#)

You need to build an app that will use Azure AI Vision to analyze and detect animals in images.

Which type of project should you use?

☒ image classification

This answer is incorrect.

☐ image detection

☐ object analysis

☐ object classification

☐ object detection

✓ This answer is correct.

Object detection returns the coordinates in an image where the applied label(s) can be found, while image classification applies one or more labels to an entire image.

[What is Custom Vision? - Azure Cognitive Services | Microsoft Learn](#)

[Detect objects in images - Training | Microsoft Learn](#)

You are creating a model to detect animals in wildlife images.

You need to train a model that identifies and locates animals efficiently.

What should you use?

- ☐ Build a CNN model by using Azure Machine Learning.
- ☐ Train a classification model by using Azure Custom Vision.
- ☒ Use Azure Custom Vision to train an object detection model.
- ☐ Use the prebuilt image tagging model in Azure AI Vision.

✓ This answer is correct.

Using Azure Custom Vision to train an object detection model is the optimal choice because it meets the requirements for identifying and locating animals in wildlife images. Building a CNN model by using Azure Machine Learning is resource-intensive and unnecessary given the availability of Azure Custom Vision. Training a classification model by using Azure Custom Vision does not provide localization information, which is essential for this task. Using the prebuilt image tagging model in Azure AI Vision lacks customization and localization capabilities, making it unsuitable for this scenario.

[Enrich a search index in Azure AI Search with custom classes and Azure AI Language - Training | Microsoft Learn](#)
[Get started with Azure Document Intelligence - Training | Microsoft Learn](#)

Your organization has trained a custom vision model using Microsoft Azure Custom Vision to classify product images. The model has satisfactory performance metrics.

You need to make the model available for predictions in a production environment.

What should you do next?

- ☐ Create a new Azure resource for the model.
- ☐ Export the model for local use.
- ☒ Publish the model and retrieve its Prediction URL and key.
- ☐ Re-train the model with additional data.

✓ This answer is correct.

Publishing the trained iteration and retrieving its Prediction URL and key is the appropriate action to make the model accessible for programmatic predictions in a production environment. Exporting the model for local use is not aligned with the goal of leveraging the Microsoft Azure Custom Vision Prediction API. Re-training the model is unnecessary as the current performance metrics are satisfactory and does not contribute to the deployment process. Creating a new Azure resource is unrelated to the task of deploying the trained model and does not address the stated goal.

[Get started with Azure Document Intelligence - Training | Microsoft Learn](#)

[Create a client application that uses an image generation model - Training | Microsoft Learn](#)

You are developing a custom vision model using Microsoft Azure AI Custom Vision to detect specific objects in images. The dataset includes labeled images for training and validation.

You need to ensure that the trained model meets the required accuracy by assessing its performance.

What action should you take?

- ☐ Deploy the model without testing.
- ☒ Evaluate precision and recall metrics.
✓ This answer is correct.
- ☐ Source new images for testing.
This answer is incorrect.
- ☐ Retrain the model with a new dataset.

Evaluating precision and recall metrics is the most effective way to assess the performance of an object detection model, as these metrics provide a detailed understanding of its accuracy and ability to identify objects correctly. Deploying the model without testing bypasses the critical step of ensuring its reliability and effectiveness. Sourcing new images for testing isn't necessary with a dataset already containing validation images. Retraining the model with a new dataset focuses on improving data quality rather than assessing the current model's performance, making it unsuitable for the given objective.

[Test and retrain a Custom Vision model - Training | Microsoft Learn](#)

You are building an app that will use the Azure AI Custom Vision API to detect when all the spaces in a parking lot are empty.

Which feature of the API should you use?

☐ image classification

☐ image description

☒ object detection

✓ This answer is correct.

Object detection is similar to image classification, but it returns the coordinates in an image where the applied label(s) can be found.

Image description analyzes an image and generates a human-readable phrase that describes its contents. Image classification applies one or more labels to an entire image.

[What is Custom Vision? - Azure Cognitive Services | Microsoft Learn](#)

[Detect objects in images - Training | Microsoft Learn](#)

You are building a video processing app that will use Azure AI Video Indexer.

You need to configure the training and learning phases for the app. The solution must train the model based on the probability of specific word combinations by using a custom Language model.

Which three practices should be followed for the training data? Each correct answer presents a complete solution.

- ☐ Include at least 500,000 sentences.
- ☒ Include multiple examples of spoken sentences.
✓ This answer is correct.
- ☐ Include special characters such as ~, #, @, %, and &.
- ☒ Provide multiple adaptation options.
✓ This answer is correct.
- ☒ Put only one sentence per line.
✓ This answer is correct.
- ☐ Repeat the identical sentence multiple times.

When training the model, you should avoid repeating an identical sentence multiple times, as it may create bias against the rest of the input.

You should avoid including uncommon symbols (~, # @ % &), as they will be discarded. The sentences in which they appear will also be discarded.

You should also avoid putting inputs that are too large, such as hundreds of thousands of sentences, because doing so will dilute the effect of boosting.

[Customize a Language model in Azure Video Indexer - Azure - Azure Video Indexer | Microsoft Learn](#)

[Analyze video - Training | Microsoft Learn](#)

You are building an app that will analyze the sentiment of user feedback by using Azure AI Language.

You have a test document named Test.docx that contains one positive sentence and multiple neutral sentences.

You need to validate the app by analyzing Test.docx.

Which label will the app return for Test.docx?

☐ mixed

☐ negative

☒ neutral

This answer is incorrect.

☐ positive

✓ This answer is correct.

If there is at least one positive sentence in the document, and the rest of the document is neutral, then the document label is positive.

[How to perform sentiment analysis and opinion mining - Azure Cognitive Services | Microsoft Learn](#)

[Extract insights from text with the Language service - Training | Microsoft Learn](#)

You are building an Azure AI Translator custom model.

You need to ensure that the translation accuracy for the model has a Bilingual Evaluation Understudy (BLEU) score that indicates high quality.

What is the minimum score range required?

☐ 0 to 19

☐ 20 to 39

☒ 40 to 59

✓This answer is correct.

☐ 60 to 79

☒ 80 to 100

This answer is incorrect.

between 40 and 60 indicates a high-quality translation..

[Key terms - Custom Translator - Azure Cognitive Services | Microsoft Learn](#)

[Translate text with the Translator service - Training | Microsoft Learn](#)

You are building an app that will recognize the intent and entities of user utterances in real-time.

You need to implement the pattern matching intent recognition mechanism. The solution must only detect entities that you define in a catalog of phrases.

Which entity type should you use?

☐ the List entity using Fuzzy mode

☒ the List entity using Strict mode

✓ This answer is correct.

☐ the prebuilt entity using Fuzzy mode

☐ the prebuilt entity using Strict mode

☐ the regex entity using Fuzzy mode

☐ the regex entity using Strict mode

The List entity is made up of a list of phrases that will guide the engine on how to match the text. When an entity has an ID of type List and is in Strict mode, the engine will only match if the text in the slot appears in the list.

[Pattern Matching overview - Azure Cognitive Services | Microsoft Learn](#)

[Create speech-enabled apps with the Speech service - Training | Microsoft Learn](#)

You are building an app that will recognize the intent and entities of user utterances in real-time.

You are evaluating the use of intent recognition with the Azure AI Speech and Azure AI Language services or simple pattern matching.

When should you use pattern matching?

☒ You are only interested in matching strictly what the user said.

✓ This answer is correct.

☐ You must manage the model by using a web app.

☐ You must use a machine learned entity.

☐ You must use a prebuilt entity.

The only option is to use pattern matching over Language Understanding when you want to strictly match what a user said, as the incorrect options are only available in Language Understanding.

[How to recognize intents with custom entity pattern matching - Azure Cognitive Services | Microsoft Learn](#)

[Create speech-enabled apps with the Speech service - Training | Microsoft Learn](#)

Question 47 of 50

You are building a model that uses Conversational Language Understanding (CLU).

You need to train the model.

Which training methods can you use?

☒ advanced, deterministic, and standard

This answer is incorrect.

☐ advanced only

☐ deterministic only

☒ standard and advanced only

✓ This answer is correct.

☐ standard only

Both standard and advanced are from CLU. Deterministic is a method from Language Understanding.

[How to train and evaluate models in Conversational Language Understanding - Azure Cognitive Services | Microsoft Learn](#)

[How to use train and test - Azure Cognitive Services | Microsoft Learn](#)

[Build a Language Understanding model - Training | Microsoft Learn](#)

Question 48 of 50

You are building a multilingual conversational app by using Conversational Language Understanding (CLU), part of Azure AI Language service.

You create a CLU model that will serve multiple languages.

You need to optimize the performance of the model. The solution must minimize development effort.

What should you do?

- ☒ Add utterances for languages that are performing poorly in the model.
✓This answer is correct.
- ☐ Configure the app to only query utterances in the language that was used to train the model.
- ☐ Create separate projects for each language.
- ☐ Train the model by using utterances in multiple languages and only query the model by using the project language.

This answer is incorrect.

With CLU, there is no need to use multiple projects for a model. For example, you can train a model in English and query it in German. There is no project language, therefore, adding utterances for languages in the model that are performing poorly is the appropriate solution to increase performance.

[Multilingual projects - Azure Cognitive Services | Microsoft Learn](#)

[Build a Language Understanding model - Training | Microsoft Learn](#)

You plan to build a chatbot that will help users answer FAQs.

You need to identify which scenarios are suitable for use with the Azure AI Language question answering service.

Which three scenarios should you identify? Each correct answer presents a complete solution.

☐ when you have a bot conversation that includes dynamic information

☒ when you have a bot conversation that includes static information

✓ This answer is correct.

☐ when you have dynamic information in a knowledge base of answers

☒ when you have static information in a knowledge base of answers

✓ This answer is correct.

☐ when you need to provide the same answer to a request, question, or command

✓ This answer is correct.

☒ when you need to provide unique answers to each request, question, or command

This answer is incorrect.

Question answering only works with static information, not with dynamic information. In addition, it will always provide the same answer to the same question.

[What is question answering? - Azure Cognitive Services | Microsoft Learn](#)

[Build a question answering solution - Training | Microsoft Learn](#)

You are building a chatbot that will use the Azure AI Language question answering service.

You need to identify the operational costs for the service.

Which two parameters will influence the costs? Each correct answer presents a complete solution.

☐ the number of assigned metadata tags

☐ the number of knowledge base editors

☐ the number of supported languages

☒ the required throughput

✓ This answer is correct.

☒ the size and the number of knowledge bases

✓ This answer is correct.

The throughput, the size, and the number of knowledge bases affect the pricing tier, whereas the other parameters do not affect pricing.

[Azure resources – question answering – Azure Cognitive Services | Microsoft Learn](#)

[Build a question answering solution – Training | Microsoft Learn](#)